

TAPAS: Generating Parallel Accelerators from Parallel Programs

<https://github.com/sfu-arch/tapas>

Steven Margerm, Amirali Sharifian, Apala Guha, Arrvinth Shriraman
School of Computing Science
Simon Fraser University
{smargerm, amiralis, aguha, ashriram}@cs.sfu.ca

Gilles Pokam
Intel Corporation
{gilles.a.pokam}@intel.com

Abstract—High-level-synthesis (HLS) tools generate accelerators from software programs to ease the task of building hardware. Unfortunately, current HLS tools have limited support for concurrency, which impacts the speedup achievable with the generated accelerator. Current approaches only target fixed static patterns (e.g., pipeline, data-parallel kernels). This constraints the ability of software programmers to express concurrency. Moreover, the generated accelerator loses a key benefit of parallel hardware, dynamic asynchrony, and the potential to hide long latency and cache misses.

We have developed *TAPAS*, an HLS toolchain for generating parallel accelerators from programs with dynamic parallelism. *TAPAS* is built on top of Tapir [22], [39], which embeds fork-join parallelism into the compiler’s intermediate-representation. *TAPAS* leverages the compiler IR to identify parallelism and synthesizes the hardware logic. *TAPAS* provides first-class architecture support for spawning, coordinating and synchronizing tasks during accelerator execution. We demonstrate *TAPAS* can generate accelerators for concurrent programs with heterogeneous, nested and recursive parallelism. Our evaluation on Intel-Altera DE1-SoC and Arria-10 boards demonstrates that *TAPAS* generated accelerators achieve 20× the power efficiency of an Intel Xeon, while maintaining comparable performance. We also show that *TAPAS* enables lightweight tasks that can be spawned in ≈ 10 cycles and enables accelerators to exploit available fine-grain parallelism. *TAPAS* is a complete HLS toolchain for synthesizing parallel programs to accelerators and is open-sourced.

Index Terms—High-level Synthesis, LLVM, Chisel, HLS, Cilk, TAPAS, Hardware accelerator, Power efficiency, Dynamic parallelism, FPGA

I. Introduction

Industry and academia realize that hardware customization is required to continue performance scaling as semiconductor scaling tapers off. Amazon EC2 [23] and Huawei have made FPGAs available to the public through the cloud. Microsoft [35] is also exploiting FPGAs for accelerating datacenter services. To address the challenges of developing application or domain specific hardware, high-level-synthesis (HLS) tools have been introduced. HLS translates a program in high-level language (e.g., C, C++) to an RTL circuit specification. It is an open question whether HLS tools have enough flexibility to permit software engineers to design high performance hardware.

A key limitation of HLS tools is their approach to concurrency. Accelerators attain high performance by instantiating multiple execution units that effectively support both coarse-grain and fine-grain concurrency [1], [5], [34] (relative to software). Unfortunately, current HLS tools do not effectively support concurrent languages. HLS tools also require an extensive set of annotations to generate parallel architectures. concurrency [10]. High-level-synthesis (HLS) tools with C interface typically analyze loops and employ techniques such as unrolling and pipelining [1]. Both Intel and Xilinx have targeted HLS at data parallelism [26].

HLS tools were aware of the challenges introduced by concurrency and have sought to exploit higher-level parallelism. LegUp [2], [11] includes support for a subset of the OpenMP and pthread APIs, and seeks to benefit from thread-level parallelism. IBM’s liquid metal [5] supported streaming kernel parallelism. Both toolchains are limited to static concurrency patterns, i.e. the parallelism structures are known during hardware generation and the structures cannot change during execution.

Recent works and industry-standard HLS tools have adopted fixed hardware templates that target specific concurrency patterns. Common templates include data parallelism, loop parallelism and loop pipelining [29], [33], [42]. The application programmer is expected to annotate and modify the application to fit the template. Template-based HLS adopts a construct-and-run approach in which the concurrency and operations are scheduled statically at hardware generation time. Unfortunately, in many concurrent programs the parallelism evolves as the program runs, either due to control flow [30], or run time non-determinism [14], [15] (see example in Figure 1).

Current HLS tools are built on a sequential compiler i.e., compiler intermediate representation and passes restricted to a sequential program-dependence-graph. Hence, prior tools largely focused on programs with static parallelism that can be expressed through templates (e.g., pragma pipeline) or library calls (e.g., OpenCL). Our work focuses on programs with irregular fine-grain parallelism expressed implicitly within the program and it has been built on a parallel compiler released in 2017 [39]. We demonstrate that for programs with dynamic concurrency, FPGAs can achieve higher performance/watt than a multicore.

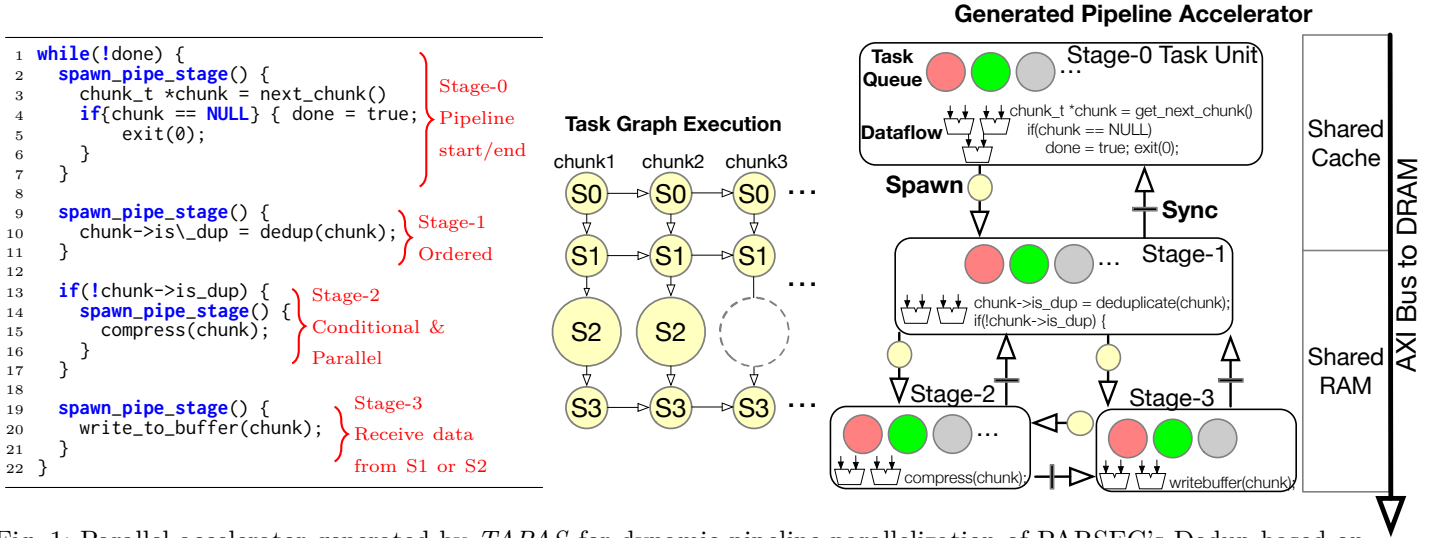


Fig. 1: Parallel accelerator generated by *TAPAS* for dynamic pipeline parallelization of PARSEC’s Dedup based on Cilk-P [28] (modified to enhance clarity).

Our Approach

Our work focuses on synthesis of hardware accelerators from parallel programs that contain on-the-fly or dynamic parallelism. *TAPAS* is a complete HLS framework that leverages parallel IR [41] to generate the RTL for a parallel task-based accelerator. The accelerator architecture includes support for spawning and synchronizing both homogeneous and heterogeneous tasks at run time. *TAPAS* leverages the parallelism markers embedded by Tapir to generate RTL in two stages. The first-stage analyzes the parallel IR to infer the task dependencies, required synchronization, and generates a top-level architecture at the granularity of tasks. In the second stage, it generates the dataflow execution logic for each task; we permit arbitrary control flow (including loops) and memory operations. The microarchitecture generated by *TAPAS* is specified in parameterized Chisel [4] and permits the designer to vary the number of tiles dedicated-per task, resource per task (e.g., queue depth, registers, scratchpad) and memory system capacity.

We illustrate that the dynamic task-based accelerator has flexibility for realizing nested, heterogeneous, recursive, irregular or regular concurrency patterns. We briefly discuss how *TAPAS* handles the challenges of generating hardware for a dynamically pipelined program, Dedup from PARSEC (see Figure 1); the figure includes the commented pseudo code. HLS tools find this particular code sample challenging and cannot generate an optimal microarchitecture. First, the stages in the pipeline change based on the inputs. As shown in the task graph, for some iterations stage-2 could be entirely skipped based on the results from stage-1. Second, the stages have different ordering constraints and exhibit nested parallelism. Stage-2 is embarrassingly parallel while stage-1 enforces ordering across each sequence. Finally, the pipeline termination condition (see line 4) needs to be evaluated at runtime and cannot be

statically determined (e.g., bounded loop).

To handle dynamic parallel patterns *TAPAS* generates a hierarchical microarchitecture that includes first-class support for generic tasks. At the top-level the accelerator’s microarchitecture consists of a collection of unique atomic task units (one for-each heterogeneous task in the system). Each task unit internally manages the dataflow logic for executing the task. The generated architecture has the following benefits: i) dynamic task spawning enables the program control to skip stages entirely and change the pipeline communication pattern, ii) the hierarchical task logic organization permits concurrent tasks to be nested. *TAPAS* permits Dedup’s stage-2 to be internally parallelized while ordering the tasks for stage-1. iii) The architecture eliminates dedicated communication ports, and allocates local RAM for communicating data between the tasks. This permits Dedup’s stage-1 to directly pass data to stage-3 when stage-2 is bypassed conditionally. iv) Finally, the architecture does not require any separate control for managing the task dependencies. *TAPAS* derives the concurrency control from the compiler IR and embeds it within the tasks e.g., pipeline exit function is the `next_chunk()` dataflow embedded within stage-0.

- 1) We have developed *TAPAS*, a complete open-source HLS tool that generates parallel hardware accelerators with support for dynamic task parallelism.
- 2) *TAPAS*’s framework is based on a parallel compiler intermediate-representation and includes support for arbitrarily nested parallelism and irregular task parallelism. It is language agnostic and has been tested using Cilk, Cilk-P and OpenMP.
- 3) We have developed a library of hardware components for spawning and synchronizing tasks, buffering tasks, and inter-task communication. We demonstrate that *TAPAS* HLS can compose these components to generate high performance parallel accelerators.

- 4) We evaluate the performance and flexibility of *TAPAS* on the Intel-Altera DE1-SoC and Arria 10 FPGA boards. *TAPAS* achieves $\simeq 20\times$ better performance/watt than an Intel Xeon quad core, while the performance is comparable to the multicore processor.

II. Background and Scope

There is a gap in the quality of the hardware generated by HLS and human-designs a) partly due to the inability of the HLS tool to comprehend and exploit the parallelism available in software, and b) partly due to underlying abstractions being not available in the hardware architecture. As resource abundant FPGAs appear in the market it becomes imperative for HLS to support dynamic parallelism where the user only specifies what tasks can run in parallel, instead of how the parallel tasks are mapped to execution units.

A. Why dynamic task parallelism in *TAPAS* ?

A question that naturally arises when generating a parallel architecture is how does one specify parallelism to an HLS tool? There appears to be no consensus among current toolchains. This is primarily a result of there not being a standard framework to support concurrent execution, as is true for CPUs. Common frameworks such as OpenMP [2], [7] and Intel TBB are implemented using threads. However, it is unclear if the requirements of threads (e.g., precise register context, shared memory, per-thread stack) can be supported on non-CPU architectures at low overhead. Recent works have included support for threads in OpenMP loops [3], [41], [43].

We present an alternate vision based on the task abstraction. Please note that the notion of dynamic tasks [3], [22] we discuss here is different from the notion of static tasks explored in prior work [5]. Intuitively, a task is analogous to some encapsulated computation in software (not unlike a function call) which takes arguments and produces a value after running to completion. Every task is described as a three tuple ($f()$, args, sync). The function $f()$ represents a scoped subset of the program dependence graph which implements the functionality. Args[] is a set of arguments passed to the function and sync is a run time field that represents other tasks that need to be synchronized. The key feature of tasks is that tasks can dynamically spawn new child tasks at run time. There has been extensive work in supporting static task parallelism in FPGAs [6], [9], [10], [21]. Prior works statically scheduled these tasks on the underlying execution units and relied on the task abstraction for understanding the static concurrency pattern. *TAPAS* provides direct support in hardware for creating and synchronizing tasks dynamically based on accelerator execution. *TAPAS* adopts a dynamic software task model that has been previously explored in the context of [24], [27], [38], vector architectures [40] and GPUs [32].

B. Tapas vs Industry-standard HLS

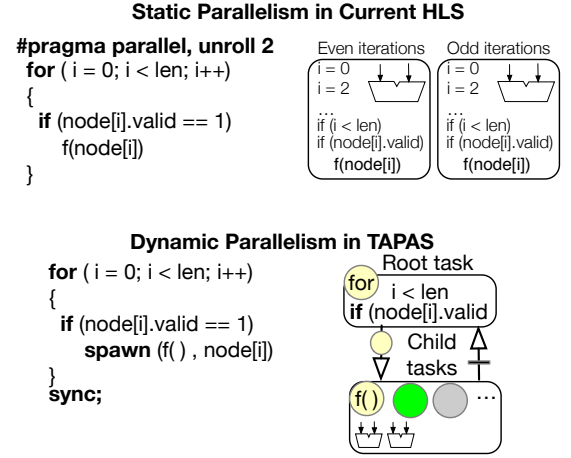


Fig. 2: Generating parallel for loop. HLS and static scheduling vs. *TAPAS* and dynamic scheduling

Dynamic Parallelism vs Static Parallelism

The term dynamic parallelism refers to enabling tasks to vary at run-time, both in terms of the type of child task spawned and the number of child tasks spawned. Prior HLS tools adopt static parallelism in which any concurrent thread/task is created and scheduled up-front. We use a commonly used parallel for-loop (Figure 2) to illustrate the differences between dynamic and static parallelism. The loop exhibits dynamic parallelism. First, the loop bounds are determined by a parameter len which is known only during execution and which varies the number of parallel loop iterations. Second, the $bar()$ function is invoked only if $node[i]$ is valid. The figure shows how *TAPAS* handles this pattern. *TAPAS* creates a root task for the loop control, which spawns a child task, $f()$, only when required (i.e., node is valid) and up to the dynamic maximum of len . Current HLS tools will unroll the loop to exploit the static parallelism. When a loop is unrolled, the HLS tools create multiple hardware execution units onto which successive loop iterations are statically scheduled at the hardware construction time (in Figure 2 unroll factor is 2). Therefore, they must plan for the worst case and allocate resources for all possible iterations regardless of whether they are actually executed, and must handle corner cases (e.g. $len \neq 0$ unroll).

Static vs Dynamic Scheduling

Another limitation of current industry-standard HLS tools [16] is the lack of dynamic scheduling i.e., the ability for the dataflow to handle variances in instruction latency (e.g., cache misses). Even the HLS tools that support threads [12], only support static scheduling of instructions. Since memory instructions also need to have deterministic scheduling, prior HLS tools primarily support a streaming memory model in which data is loaded into a scratchpad ahead of invocation. While the combination of static scheduling, static concurrency, and streaming memory model leads to high efficiency, but it limits the type of

workloads that HLS can target. A pre-requisite for supporting dynamic task parallelism is shared memory and caches. Consequentially, *TAPAS* needs to handle non-deterministic latency in memory operations (see Section III-C).

Concurrent work from Josipović et al. [25] at FPGA 2018 has started investigating the benefits of dynamic scheduling of instructions. However, their work only exploits static parallelism from loops in sequential programs. *TAPAS*'s focuses on dynamic parallelism and parallel programs. *TAPAS* includes support for the task abstraction in the compiler that makes it feasible to target parallel languages such as Cilk. *TAPAS* also supports dynamic scheduling, however this is not our focus.

C. *TAPAS* vs Prior Work

TABLE I: Comparing the features of HLS tools

	Base HLS [1], [8]	HLS+Kernels [2], [10], [19], [26], [42]	Pattern [18], [29], [31], [34]	TAPAS
Scope	Seq.	Kernel	Pattern	Parallel Prog.
Target	Loops	Thread	Pattern	Tasks
Hints	#pragma	Kernels	Pattern	Spawn/Sync
Dynamic	—			✓
Loops	Unroll/Pipeline		Parallelize	
Heterogeneity	—	Limited		✓
Nested parallel	—		Limited	✓
Lightweight	✓	—		✓
Multithreading	—			✓
Multicore	—	✓ Multiple execution units		

Table I summarizes the feature set of current HLS tools. Many HLS tools [16] primarily target sequential programs and unroll loops to exploit instruction parallelism. A parallel architecture is often realized by using the HLS compiler to synthesize a single hardware core, and then typically requires an expert to manually instantiate multiple instances of the core in a hardware description language. To avoid this, both Xilinx Vivado and Intel HLS unroll and pipeline loops to convert loop parallelism to instruction level parallelism. Achieving efficient hardware requires the software developer to identify where it might be feasible to exploit loop parallelism and add additional hardware-oriented pragmas. HLS tools have anticipated the need to target higher levels of parallelism. Recent works have supported a subset of OpenCL, OpenMP or Pthreads. The primarily target is data parallel kernels. Current HLS tools schedule the concurrent operations statically and do not support dynamic spawning, asynchronous behavior, or nested parallelism. Furthermore, since the HLS tools statically schedule the memory operations, they require code annotations to help identify the streaming and FIFO access patterns between functions [13]. Finally, a promising

avenue of research is HLS for domain-specific patterns. The hardware expert designs a parameterized template that targets a parallelism pattern (e.g., pipeline) and the software developer modifies the applications to ensure the program structure matches the pattern. Unfortunately, patterns always risk becoming obsolete.

TAPAS targets parallel programs (not a particular pattern) and only requires the programmer to identify concurrent tasks. It is built on a parallel compiler and leverages the information to automatically synthesize parallel hardware for arbitrary task graphs. The key novelty of *TAPAS* is that tasks can dynamically spawn and sync with other tasks. This enables *TAPAS* to handle a variety of common programming patterns including nested, recursive and heterogeneous parallelism. Finally, *TAPAS* supports dynamic scheduling of operations and handles non-determinism to enable a cache-based memory model.

III. *TAPAS*: High-Level-Synthesizing Dynamic Parallel Accelerators

TAPAS is a hierarchical HLS toolchain for generating the RTL for a parallel application-specific accelerator (see Figure 3). *TAPAS* is language agnostic since it relies on Tapir-LLVM to parse the parallel program and generate compiler IR with additional markers indicating the parallelism. The input to *TAPAS* is a parallel program with markers for tasks and parallel loops; currently our infrastructure has been tested using Cilk, OpenMP and Cilk-P. Figure 3 shows the stages in *TAPAS* and RTL generation for a program with nested parallel loops. *TAPAS* consists of three stages. In Stage-1, (Section III-A) *TAPAS* analyzes the compiler IR, extracts the task dependencies, and generates the top-level RTL. The task units are declared and wired to the memory system. In Stage-2(Section III-C) the program graph of each task is analyzed, and the RTL is generated for the dataflow of each task unit. Finally, in Stage-3 we configure and set the hardware parameters (e.g., number of execution cores) based on a specific deployment (e.g., LUTs available on the FPGA) and generate FPGA bitstream.

TAPAS-generated accelerators support dynamic parallelism, dynamic scheduling, and caches. We restrict all communication between the ARM and the accelerator to occur through shared memory. Currently, *TAPAS* maps the accelerators to the FPGA on an SoC board. The ARM and the FPGA share a 512KB L2 cache. We synthesize a 16K L1 cache for the accelerator which is kept coherent with the L2 through AXI. *TAPAS* generates a binary for the program regions/functions that cannot be offloaded (e.g., due to system calls) and they run on the ARM. *TAPAS* does not rely on any hard logic in the FPGA and synthesizes the logic required to support the parallelism. This enables a flexible execution model that is independent of the processor and enables *TAPAS* to target different FPGA boards.

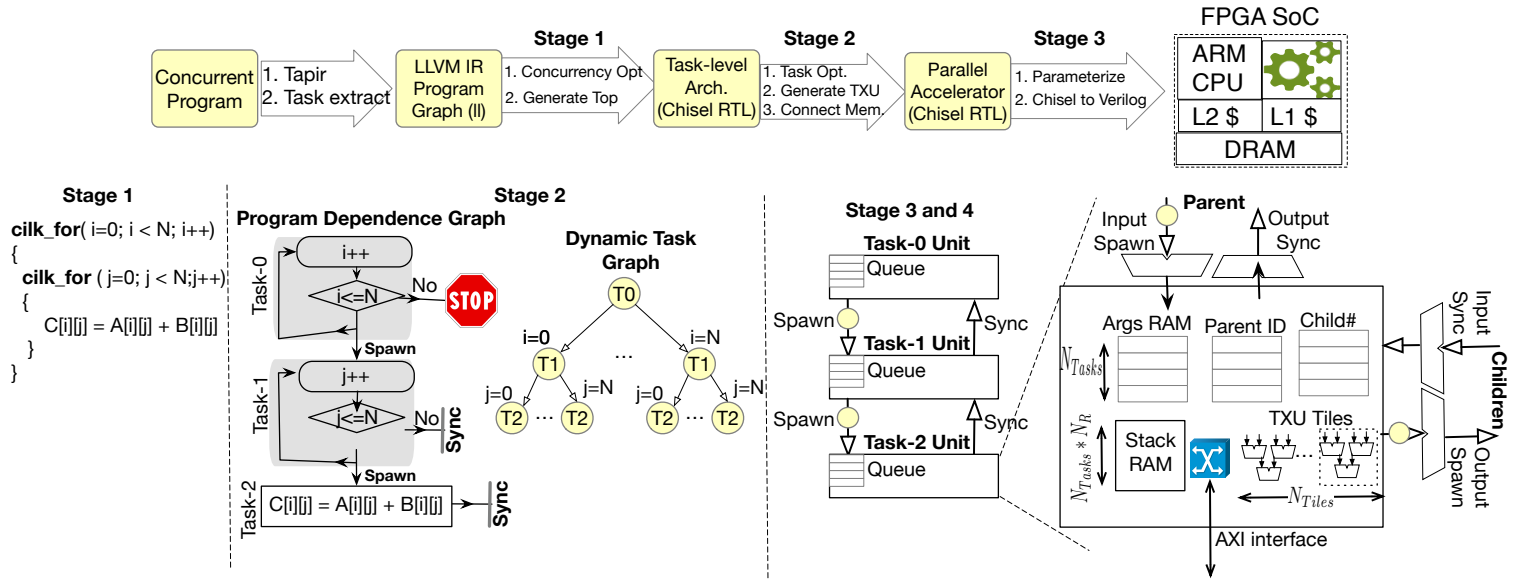


Fig. 3: Overview of *TAPAS*. Top: Compilation flow. Bottom: Generated output at each stage.

A. Stage 1: Task Parallel Architecture

TAPAS relies on Tapir [39] to comprehend the semantics required by the task-based accelerator architecture. Tapir adds three instructions to LLVM IR, *detach*, *reattach* and *sync*, to express fork-join parallel programs. Using these three instructions TAPIR can support dynamic task spawning (create a concurrent task) and sync (synchronize parent and child). We describe the front-end task compiler pass in more detail in §III-F and focus on the hardware generation itself once the task dependencies are known. The generated accelerator consists of multiple task units at the top-level, and each task unit represents a unique task. Figure 4 illustrates the top-level RTL, the interface and the parameters associated with the interface. *TAPAS* supports time multiplexing (equivalent to simultaneous multithreading) of multiple tasks on an execution unit, dynamic tiling and assignment of tasks at runtime to different execution units (equivalent to multicore). A task unit is an execution engine for a single task and serves as the basic building block in the architecture. The accelerator can consist of any number of task units interacting to create different task graphs. There are four main components within each atomic task unit: i) The task queue which manages spawned tasks, ii) Parent task interface (Spawn/Synchronization ports), iii) Child task interface (Spawn/Synchronization ports), and iv) Task Execution Unit (TXU) which represents a pipelined dataflow execution unit.

B. Execution Example

We describe the functionality of each of the components in the task unit by considering the implementation of the nested loop example (see Figure 5). The task graph in the figure illustrates three tasks **T0**, the outer loop control and spawner of N instances of inner loop. **T1** is the inner loop control and spawner of N instances of **T2**, the body. Finally **T2** performs the actual work, reading elements from

```

1 class NestedAccelerator(implicit p: Parameters)
2   extends Module {
3     // Simple L1 Cache
4     val SharedL1cache = Module(new Cache)
5     // DRAM. AXI4 interface
6     val DRAM = Module(new NastiMemSlave)
7     // Wire DRAM and L1cache to AXI
8
9     // Initialize task units
10    val Task-0 = Module(
11      new Task0(Nt=32, Nr=32, ParentPort = 1,
12        ChildPort=1, new T0-DF())
13    )
14    val Task-1 = Module(
15      new Task1(Nt=32, Nr=32, ParentPort = 1,
16        ChildPort=1, new T1-DF())
17    )
18    val Task-2 = Module(
19      new Task2(Nt=32, Nr=32, ParentPort = 1,
20        ChildPort=1, new T2-DF())
21    )
22    //***** Connect Task Units *****/
23    Task-0.io.in <-> io.in
24    io.out <-> Task-0.io.out
25    Task-1.io.detach.in <-> Task-0.io.spawn.out
26    Task-0.io.sync.in <-> Task-0.io.out
27    Task-2.io.detach.in <-> Task-1.io.spawn.out
28    Task-2.io.sync.in <-> Task-1.io.out
29    //***** Connect Cache to Task units*****/
30  }

```

Fig. 4: *TAPAS* generated microarchitecture in Chisel [4].

the $A[i][j]$, $B[i][j]$ and adding them. In this application, N dynamic instances of task **T1** will be created (for each iteration of the outer loop) and each dynamic instance of **T1** will create N instances of **T2** (total: N^2 instances).

A task in the queue can be in of the following states

- **READY**: spawned, but not allocated a TXU
- **EXE**: TXU allocated, but task has not complete
- **COMPLETE**: execution complete and need to synchronize with the parent
- **SYNC**: Waiting on synchronization with child tasks. The task queue metadata consists of a child join counter (**Child#**), the **ParentID** and **Args[]** RAM (argument RAM).

① illustrates a spawn operation, with **T0** initiating the task corresponding to the inner loop iteration, **T1**. A spawn is a tuple, **Args[]** and **ParentID**. The **ParentID** consists of

Task Status: R Ready E Active C Complete S Sync waiting

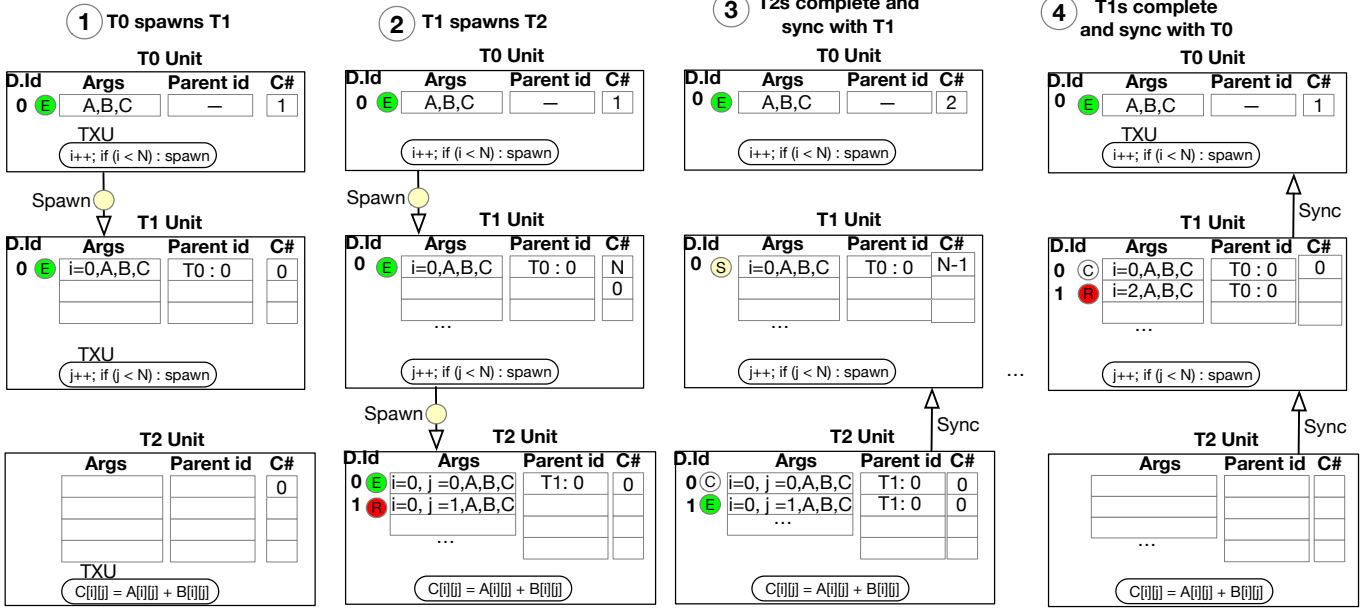


Fig. 5: Execution flow of Nested-loop accelerator generated by TAPAS

[SID, DyID]. The SID refers to the name of the parent task (in this instance T0) and the DyID corresponds to the task queue entry allocated to the instance of the parent task (index 0 here). This corresponds to dynamic task T0:0 spawning an instance of T1 (a j-loop, when $i = 0$). The ParentID metadata available in the spawn is noted down in the allocated to the spawned T1 task and will be used during synchronization. In (2) the dynamic instance T1:0 (corresponding to the inner j loop with $i = 0$) creates N instances of the inner body T2. The field C# (Child#), in task T1:0's queue entry corresponds to the count of the children tasks that are created by dynamic task T1:0. In this example, N instances of T2 are created corresponding to loop iterations $i=0, j=0..N-1$. Note that, task T0 may concurrently create other instances T1:1, T1:2, ... (inner j loop for $i=1, i=2, \dots$ iteration) if there is enough queuing available. The task unit asynchronously assigns task execution units for the ready tasks.

In (3) as the instances of T2 complete they synchronize with their parent task that spawned them. Each task will only synchronize and join with the parent task that created it. Here, the T2 instances T2:0...T2:N-1 (corresponding to tasks $i=0, j=0..N-1$) will join on completion with the dynamic instance of their parent T1:0 ($i=0, j$ -loop control). Joining entails decrementing the counter in the queue entry (index 0) corresponding to T1:0. The purpose of noting down the SID and DyID when the T2 tasks were spawned is clear now. The SID permits composability and allows heterogeneous task units to communicate with and dynamically spawn a shared task. The SID serves as the network id of the parent task unit to route back on a join. The DyID serves as the index into the queue

within the task unit corresponding to the SID. Finally in (4), once T1:0 has joined with all its spawned children, it proceeds to move from SYNC to COMPLETE status and reattaches back with its parent, T0. The task queue interfaces decouple task creation from task execution. The spawn and sync are asynchronous, and employ ready-valid signals. The asynchronous design permits us to vary the resource parameters per task without having to reschedule the tasks to deal with changes in latency.

C. Stage 2: Generating Task Exe Unit (TXU)

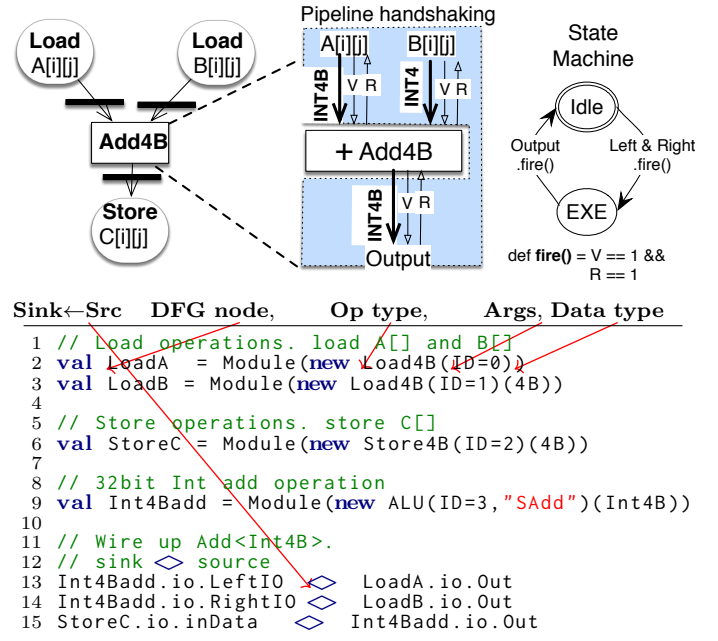


Fig. 6: Task-2's Task Execution Unit (TXU).

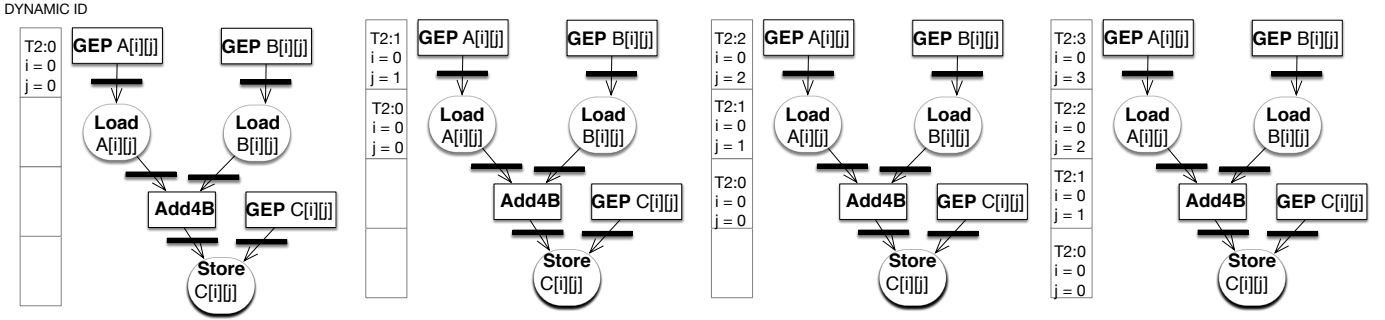


Fig. 7: Multiple tasks simultaneously outstanding on TXU.

TXU is the representation of execution engine within each task unit. Each TXU is a fully pipeline execution unit which permits multiple dynamic instances of a task execute simultaneously. The TXUs only communicate at the task boundaries with each other. All the inter-TXU communication is marshaled through a shared scratchpad or the cache.

TAPAS generates the logic for the TXUs based on the per-task sub-program-dependence-graph earmarked by our compiler. Each TXU is a dataflow that enables fine-grain instruction level parallelism to be mined. *TAPAS* HLS dynamically schedules the operations in the TXU. An automatic pipelining process introduces latency insensitive ready-valid interfaces between each operation in the dataflow. A dataflow graph mapped to the TXU may contain nodes with multi-cycle latency (e.g. floating point operations), and non-deterministic latency (e.g., memory operations). This approach is in contrast to current industry strength HLS tools which try to schedule the timing all operations statically; concurrent work in FPGAs has begun to analyze the potential of dynamic scheduling [25].

Figure 6 shows the add function unit from the file created by *TAPAS*, from $C[i][j] = A[i][j] + B[i][j]$. The add function unit communicates with Load A[i][j], Load B[i][j], + and Store C[i][j] via decoupled handshaking signals which contain ready and valid signals in addition to data. The handshaking interface is governed by a simple state machine. This dataflow permits multiple concurrent T2 tasks to be outstanding at the same time on the execution unit. Task pipelining is illustrated in Figure 7. The dynamic task ids correspond to the queue index allocated at run time. Note that the pipeline of a TXU is in dataflow order and tasks complete in order of issue. Any load stalls cause the pipelined dataflow to throttle and eventually stall; however this leads to a simpler implementation compared to dynamic dataflow [20].

D. Stage 3: Parameterized Accelerator

TAPAS is a parameterized hardware generator and seeks to permit late stage parameter binding. As hardware designs grow in complexity, modularity becomes necessary. The asynchrony and latency insensitivity permits each of the task units to be parameterized independently. As

shown in Figure 4 every task unit provides the mechanism for passing parameters prior to hardware elaboration and bitstream generation. While each tile has multiple parameters that can be set including the width and types of the args RAM, there are primarily two parameters that are set at this stage in toolchain, the task queue size (N_{tasks}) and the number of task execution units (N_{tiles}). We permit the user to vary the parameters on a per-task basis. The latency of the individual tasks and task dependencies will determine N_{tasks} . Determining N_{tiles} is more involved as it depends on the the processing rate required of that particular task unit and how many active tasks are required to potentially hide memory latency.

E. Task Memory interface and Memory Model

In this study, we consider a heterogeneous SoC where both processing cores and accelerator are integrated into a single chip. Each of the accelerator's caches is connected to the last-level cache, which is shared with the ARM processor over the AXI bus. A key question is how does the memory model look like on the accelerator side. *TAPAS* permits arbitrary task graph patterns to be converted into accelerators and thus needs to support a more flexible cache-like interface.




Fig. 8: Data Box. Interfaces with memory operations in logic box and transfers operations to/from a cache or scratchpad.

In *TAPAS* all the task units share an L1 cache; it is conceivable that this model is most suitable for handling a programming model most familiar to software programmers [36], [37]. Generating an optimal cache hierarchy is beyond

the scope of this paper and we primarily focus on how to route values from the cache to the TXUs.

The *data box* (see Figure 8) connects a memory operation in the TXU to a memory interface. We currently support both cache and scratchpad (we only evaluate the cache memory model in this paper). We choose to group the common logic for multiple memory operations (e.g., mis-alignment) into the data box to minimize resource requirements. Figure 8 shows the architecture of the data box. Each data box consists of the following parameterized microarchitecture components: i) an in-arbiter tree network that arbitrates amongst requests to the memory interface, ii) an out demux network that routes responses back to the memory operations in the TXU’s dataflow, and iii) a table of staging buffers that contain the actual logic for reading the required bytes from the cache/AXI memory interface (which only supports word granularity accesses). Both the request and response networks are statically routed.

F. Compiler Front-end: Tasks from IR

TAPAS is language agnostic and relies on the parallel IR introduced by Tapir [39]. Tapir provides the front-end language bindings that translates Cilk/OpenMP programs to the LLVM IR. Tapir adds three instructions to the LLVM IR, detach (or spawn), reattach and sync. Spawn and reattach together delineate a task. A detach instruction terminates the block that contains it, spawns a new task starting from the target block and continues execution in parallel from the continuation. The reattach terminates the task spawned by a preceding detach instruction. Since Tapir assumes a generic threadpool execution model, it leaves the markers in place in the original PDG (Program Dependency Graph). We leverage the markers to perform reachability analysis and extract an explicit task graph, which is the architecture blueprint for our parallel accelerator. Nesting loops and irregular flows are analyzed in this stage and in the resulting task graph all task relations and the basic blocks that constitute a task are explicitly specified. We perform live variable analysis to extract and create the requisite arguments that need to be passed between tasks; these are used to parameterize the spawn port and args RAM for each task unit.

IV. TAPAS-generated Accelerators

There are no standard benchmark suites that target only dynamic parallelism. Table II gives a brief summary of the accelerator benchmarks and shows the characteristics of each application. Our emphasis is on being able to implement accelerators for these workloads without requiring additional effort from the programmer. Here we study applications that use common software patterns that current HLS tools either find challenging and may throw errors.

A. Nested Parallel and Conditional Loops

Related benchmarks: Matrix addition, Stencil, Image scaling, Saxpy Stencil

Multiple workloads employ the similar pattern of nested

```

1 /* current: Basic block processed TaskDep: Edges between tasks
2   TFG: Task list and sub-pdg */
3 current = {BB-start}; TaskDep = {Root}; TFG = {};
4 while(current != EMPTY) {
5   startBB = current.top
6   visited[startBB] = true
7   TaskDep.top.insert(bb)
8   for-each-successor succ of bb:
9     if (visited[succ] == true) continue;
10    current.push(succ)
11    if (Edge(bb, succ) == SPAWN) {
12      child = new Task(succ)
13      TaskDep.push(child)
14      TFG.insert(child)
15      TFG.insertEdge(TaskDep.top, child, SPAWN)
16    } elseif (Edge(bb,succ) == REATTACH) {
17      Task currentTask = TaskDep.pop()
18      TFG.insertEdge(tst.top, currentTask, REATTACH)
19      TST.insert(succ)
20    } else { /*regular edge*/
21      TST.top.insert(succ)
22      dfs(TFG, current, TST)
23    }
24 }

```

Annotations in the code:
- Line 5: *Start block for*
- Line 6: *analysis*
- Line 8: *Reachable blocks from start*
- Line 11: *Nested.*
- Line 15: *Parent Spawn Child*
- Line 16: *Sync*
- Line 23: *Input.pop()*

Fig. 9: TAPAS pass for extracting tasks from Tapir

TABLE II: Benchmark Properties

Name	HLS Challenge	Memory Pattern	Per-Task	
			# Inst	# Mem
Matrix Add	Nested loops	Regular	49	21
Image Sca.	Nested, If-else loops	Regular	52	25
Saxpy	Dynamic exit loops	Regular	29	16
Stencil	Nested parallel/serial	Regular	23	16
Dedup	Task Pipeline	Irregular	180	72
Merge Sort	Recursive parallel	Regular	36	52
Fibonacci	Recursive parallel	Regular	26	19

loops. However, there are variations based on the parallelism of the loop nests, loop depth and conditional loop entry/exit. Here, we briefly discuss Stencil. Stencil is an iterative kernel (Figure 10) that updates array elements in a loop. While the loop is embarrassingly parallel, the loop bounds are variable which introduces dynamic parallelism. HLS tools parallelize this pattern in two ways: 1) flattening the inner-loops to a single level or 2) having the HLS tool only parallelize the innermost loop (which is not parallel in this case), while executing the other levels serially. In contrast, TAPAS decomposes nested loops into multiple task units. Each task unit is asynchronous and can independently configure the number of tiles to exploit all the available parallelism. TAPAS supports arbitrary nesting of loops; serial and parallel loops can be nested in any order to any depth (resources permitting). Since the task unit exposes an asynchronous spawn/sync interface, TAPAS can set up each loop independently. Every loop can execute in parallel or serial fashion based on program semantics, without requiring any changes to the outer loops.

B. Pipeline Parallelism

Dedup code and hardware accelerator are outlined in Figure 1. *Dedup* has an irregular pipeline pattern. The main challenges posed by *Dedup* are:

- **Task-Level Pipeline:** HLS tools have limited support for task-level pipelines and primarily target loop


```

1 void stencil () {
2 /* Parallel for loop */
3   cilk_for (pos = 0; pos < NROWS * NCOLS; pos +=
4     ++)) {
5     /* Serial for loop */
6     for (nr = 0; nr <= 2*NBRROWS; nr++) {
7       /* Serial for loop */
8       for (nc = 0; nc <= 2*NBRCOLS; nc++) {
9         int row = (pos/NCOLS) + nr - NBRROWS;
10        int col = (pos & (NCOLS-1)) + nc - NBRCOLS;
11        if ((row < NROWS)) {
12          if ((col < NCOLS)) {
13            ...

```




Fig. 10: Stencil Accelerator

pipelining. Dedup is parallelized using tasks that have non-trivial entry and exit logic, making it challenging to convert them to loops.

- **Conditional stages:** Emerging research [11] has sought to support functional pipelines using FIFO queues that require the program to be rewritten. Unfortunately, *dedup* also has conditional pipeline stages (see S2 in Figure 1) which FIFO queues cannot support. FIFO queues fix producer and consumer stages and cannot handle conditional pipelines.
- **Intra-stage parallelism:** The FIFO ports are ordered and this would lead to *dedup* losing parallelism in the S2 stage, which permits out-of-order chunk processing (see task-dependencies in Figure 1).
- **Pipeline control:** Finally, the pipeline termination condition is dynamically determined by an exit function (`get_next_chunk()`). HLS tools do not support dynamic exits, since they statically schedule operations.

TAPAS does not suffer from these limitations since it supports dynamic spawning/syncing of tasks and execution units are assigned at runtime. Furthermore, all tasks communicate with each other through shared memory and the parallelism is not limited by extraneous hardware structures such as FIFO.

C. Recursive Parallelism

TAPAS can effectively generate accelerators for recursively parallel programs. HLS tools have traditionally not supported recursion [11], [19] due to the lack of a program stack. Nothing precludes the addition of a stack, but it would require changes to the HLS compilation framework. Figure 11 illustrates how *TAPAS* can support recursively parallel mergesort.

In mergesort, the primary function employs a divide

and conquer strategy. It partitions an array into two halves, and recurses on each half in parallel. The parent function then waits on the children and merges the sorted halves. To implement recursion it must be possible for more than single invocation of the same function to exist at the run time. Further, the data also has to be implicitly passed via a stack. *TAPAS* achieves this through the following: i) *TAPAS* precisely captures the state needed by a recursive task from the LLVM IR and implicitly manages the stack frames in a scratchpad. ii) The task controller supports dynamic scheduling and asynchronous queuing, which permits a task to spawn itself without logic loops. iii) The task controller tracks the dynamic instances to support implicit parent-child synchronization iv) Finally, all return values from the recursion are passed through shared cache.

```

1 void mergeSort(...) {
2   if (start < end) {
3     int mid = start + ((end - start) / 2);
4     /* Spawn self to sort 1st half */
5     cilk_spawn mergeSort(list, start, mid);
6     /* Spawn self to sort 2nd half */
7     cilk_spawn mergeSort(list, mid + 1, end);
8     /* Parent waits for children */
9     cilk_sync;
10    merge(list, start, mid, end)
11  }
12 }

```




Fig. 11: Accelerator for recursive mergesort.

V. Evaluation

It is challenging to find a fair baseline since dynamic parallelism is not supported by existing HLS toolchains. Running our benchmarks on the FPGA would entail changing the algorithm and memory model, and a program re-write. Even finding a baseline CPU is challenging since the Cilk and Tapir are currently x86-only. The x86 multicore's cache hierarchy is deeper and larger than the FPGAs, which makes it challenging to understand the impact of dynamic parallelism independent of the memory system. We answer the following questions: i) Can *TAPAS* support fine-grain tasks? How fine-grain can the tasks be (§ V-A) ? ii) Is the performance improvement attributable to low task spawn latency or speeding up individual tasks with dataflow execution? (§ V-C) iii) What is the baseline performance compared to an Intel i7 quad core. (§ V-C) iv) What is the energy consumption and performance/watt benefit compared to a Intel i7 (§ V-D). In all the cases, we use the same unmodified Cilk programs. v) How does

static parallelism with prior HLS tools compare against dynamic parallelism in *TAPAS* (§ V-E)

A. Parallel Task Overhead

Result: *The overhead of spawning a task on an FPGA is significantly less than a software. This enables small, fine grain tasks to scale better.*

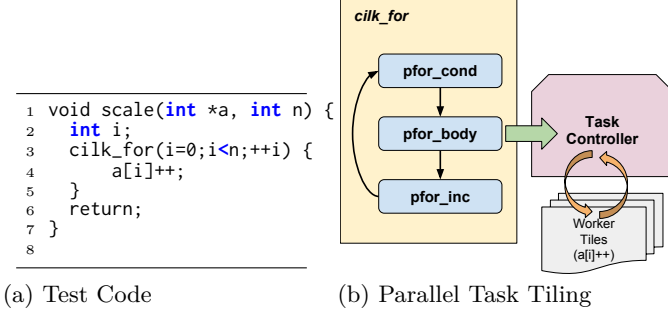


Fig. 12: Scalability Test Code

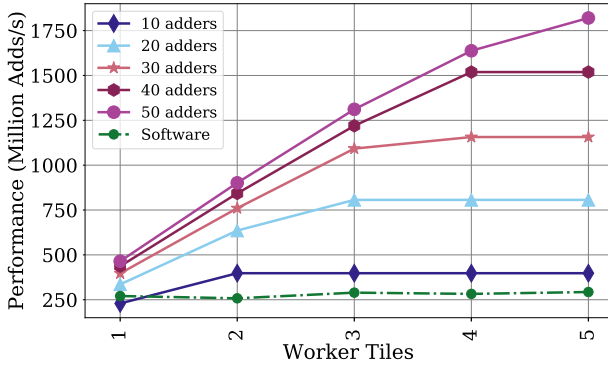


Fig. 13: Performance Scaling with Tiles

The microbenchmark in Figure 12(a) was synthesized to see how fast tasks can be spawned. Figure 12(b) provides a top-view of the generated architecture. We incrementally varied the amount of work (“+” operations) in the loop body. The performance is plotted against an increasing number of worker tiles (Figure 13). On a Arria 10 (≈ 300 Mhz) target device, we achieve a maximum spawn rate of 40 million spawns/second. Even for fine-grain tasks (50 instructions), the performance scales monotonically with the addition of parallel worker tiles. ‘Software’ in the plot (Figure 13) refers to spawning a task of 50 increments; the program was run on a Intel i7-3.4Ghz, 8MB L2 (four cores). At such fine-granularity, the software runtime for Cilk provides zero benefit due to task spawning overheads. *TAPAS* exploits the low overhead of task spawning on an FPGA and enables fine-grain parallelism not exposed in software.

B. Resource Utilization

Result: *Arria 10 FPGA board can support ≈ 100 parallel tasks each containing 100 integer operations.*

Table III shows the per-instruction and per-tile resource utilization of the accelerator. The test code was synthesized for two Intel SoC FPGAs, the Cyclone V and Arria 10 (see Table III). The primary sources of overhead in *TAPAS* are

the task controller logic and memory arbitration. On the smaller Cyclone V SoC, 10 tiles of 50 integer operations each filled 85% of the chip (153Mhz). On the larger Arria 10 board the accelerator could achieve 308 MHz and occupied 12% of the chip. A single M20K block RAM consumed is for queuing the spawned tasks in the task controller logic.

TABLE III: FPGA Utilization

MHz	Tile	Ins.	ALM	Reg	BRAM	%Chip
<i>Cyclone V (5CSEMA5)</i>						
185.46	1	1	1314	1424	1	5
178.09	1	50	2955	3523	1	10
153.61	10	1	7107	8547	1	24
159.24	10	50	24738	27604	1	85
<i>Arria 10 (10AS066)</i>						
308	10	50	28844	27659	1	12

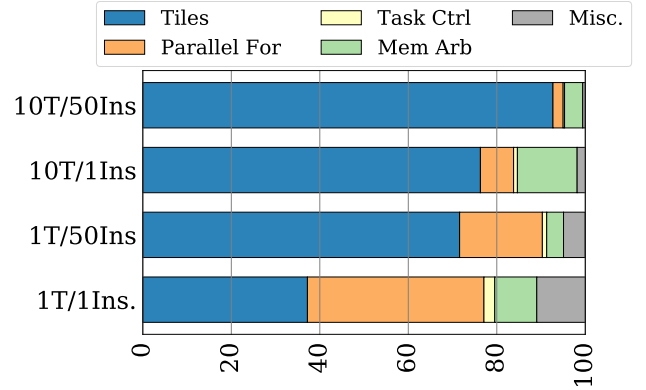


Fig. 14: ALM Utilization by Sub-block

Figure 14 shows the relative amount of ALM resources (aka. LUTs and registers) used by each sub-block of the design. In the extreme case (1 operation/task), 60% of the logic is non-compute overhead; at 50 operations/task, the overheads is $\approx 20\%$). As the number of execution tiles increases, the overhead of the control logic is amortized and at 10 tiles the control overhead is reduced to 3%. The memory network required to support shared memory access is less than 10% of overall chip resources. The network is primarily needed to support dynamic scheduling and routing values back and forth from the shared cache to the internal nodes in the task execution unit.

C. Scalability and Performance

Result 1: *TAPAS generated accelerators exploit all the available parallelism exposed by the applications and scale with increasing hardware resources (1.5–6 \times)*

Result 2: *To improve accelerator performance compared to an Intel i7 multicore a better cache hierarchy is required.*

Figure 15 plots the performance when varying the number of execution tiles per task. A performance increase from the baseline is seen in all examples with the exception of Dedup. In Dedup even the baseline case (1 tile) has four heterogeneous task units (Figure 3) organized as a pipeline, with one execution tile per task unit. Any further improvement with increasing tiles/task is feasible only if the pipeline stages are unbalanced (not the case here).

The saxpy and matrix addition improve with the addition of a second tile, but the benchmarks quickly saturate the cache bandwidth as their inner loops are small and dominated by memory reads and writes. In contrast, the Stencil benchmark is more computationally intense and consequently scales well even up to 8 tiles and beyond.

Q2. In Figure 16, we compare the execution time of accelerators against an Intel i7 quad-core (3.4Ghz, 8MB L2, 21GB/s DRAM). Identical Cilk benchmarks were used for both the i7 runs and *TAPAS*. We set the concurrency to be identical (four cores for i7 and four tiles for *TAPAS*). Accelerator designs were generated for both Cyclone V SoC FPGA and Arria 10 SoC FPGA. On the Cyclone V the accelerators performed at approximately 50% of the multicore (even achieving speedup in a few cases). On the Arria 10, the generated accelerators performed on par with the i7 as a result of higher frequency (300Mhz vs 150Mhz for the Cyclone V). The Dedup accelerator achieved best speedup since the accelerator implemented the pipeline more efficiently than software. The mergesort accelerator performed poorly in comparison to the Intel i7 since it is completely memory bound and limited by the memory system on the FPGA.




Fig. 15: Performance Scalability.

To understand the impact of the memory system, we ran the non-Cilk, sequential programs on the ARM CPU of the SoC board (same memory hierarchy as FPGA) and find it to be 13× slower than the i7. Any performance difference of our accelerators with the i7 is a result of: slower clock (FPGA 150Mhz vs 3.4Ghz i7) and smaller, less sophisticated cache hierarchy.

D. Energy Consumption

Result: *TAPAS-generated accelerators exceed the energy efficiency of the multicore often by 20×).*

In Table IV we report the absolute power consumption and resource utilization of the generated accelerators. The power is obtained using using Intel Quartus PowerPlay. It is an estimate of total power (static and dynamic) based on signal activity levels derived from gate-level simulation. The tabulated data shows how even with varied parallel patterns such as Stencil (nested loops) and Mergesort (recursive) we can effectively exploit the available resources (50% of the Cyclone V FPGA chip). Mergesort is the largest design using roughly half of the available chip resources and consuming approximately 1.5W of power. We compare the performance/watt (Figure 17) of the accelerator against a multicore; in both cases we set the concurrency level to four. The power for the multicore is directly measured through the RAPL interfaces. *TAPAS* accelerators often achieve over 20× better performance/watt than the multicore.

TABLE IV: FPGA Resources (Board: Cyclone V)

Bench	Tile	MHz	ALMs	Regs	BRAM	Power(W)
SAXPY	5	149	7195	9414	3	0.957
Stencil	3	142	11927	11543	3	1.272
Matrix	3	223	4702	7025	3	0.677
Image	4	141	4442	5814	3	0.798
Dedup	3	153	10487	6509	3	1.014
Fibonacci	4	120	5699	9887	62	1.155
Mergesort	4	134	14098	24775	74	1.491

E. Intel HLS vs *TAPAS*

An apples-to-apples comparison with prior HLS tools is challenging since: i) HLS tools only support static parallelism. It is not feasible to convert some applications to use static parallelism (e.g., recursive mergesort), and with others(e.g., Dedup) the conversion changes the algorithm entirely. ii) HLS tools deploy a streaming memory model since they statically schedule all instructions with known latencies. *TAPAS* employs caches and shared-memory, which are a pre-requisite for dynamic parallelism.

To attempt a quantitative comparison we use two benchmarks, SAXPY and Image scaling. Among our benchmarks these were amenable to static parallelism. We used the Intel HLS Compiler (v17.1) and employ the suggested streaming DRAM interface.¹ We set up the DRAM latency for both the Intel HLS and *TAPAS* to 270ns (150Mhz FPGA clock). We also set the same concurrency level. In HLS the loops was unrolled 3 times and *TAPAS* was configured to use 3 tiles. The results are listed in Table V. The results indicate that *TAPAS* is pretty competitive. It may be feasible to hand optimize the HLS implementation further, but we could also optimize *TAPAS*. The most notable difference is where the block RAMs are utilized. Intel HLS appears to generate large stream buffers in its load and store interfaces. In contrast, *TAPAS* uses a 16K L1 cache shared by all task units, but also expends block RAMs in the task queue.

¹The part_3.ddr_masters.cpp example included with Intel HLS




Fig. 16: Performance. *TAPAS* vs Intel i7.

TABLE V: Intel HLS vs *TAPAS* (Board: Cyclone V)

Bench	Tool	MHz	ALMs	Reg	BRAMs
Image	Intel HLS	155	5467	10466	67
	<i>TAPAS</i>	152	4543	7126	10
SAXPY	Intel HLS	181	3799	7961	38
	<i>TAPAS</i>	146	4254	5718	11

VI. Thoughts and Future Directions

Our results demonstrate that FPGAs and hardware accelerators have the potential to address a long outstanding challenge in concurrency, effective support for dynamic fine grain parallelism. The following facets need to be addressed to further improve performance.

- **Cache hierarchy:** To compete against a multicore processor we need to improve the overall cache hierarchy, both bandwidth and latency. The current cache macro-block we release as part of the toolchain is borrowed from the RISC-V cores with limited support for multiple outstanding cache misses. Our AXI implementation is also sub-optimal as we do not yet exploit all the burst options available in the protocol.
- **Task controllers:** The task controllers and queuing logic add latency to the critical path. In many workloads (e.g., bounded size matrix multiplication) there exist loop patterns that can be statically parallelized. *TAPAS* can benefit from statically scheduling such loops, and eliminating the task controllers. The challenge is identifying loops where this optimization may be feasible.
- **Opportunity for Dynamic Parallelism:** *TAPAS* relies on the compiler front-end (Tapir in this paper) to capture the parallelism intent of the workloads. *TAPAS* is currently capable of generating accelerators for the widely used fork-join parallelism [17]. In our current workloads, apart from the initialization, all other functions are offloaded to the accelerator.




Fig. 17: Performance/Watt: *TAPAS* vs Intel i7.

Unfortunately, our compiler front-end does not explicitly capture data-driven parallelism (e.g., inter-stage queues in a pipeline, channels in golang). As a result *TAPAS* achieves data-driven synchronization through the shared cache and memory. In the future, we plan to map data-driven parallelism to explicit hardware structures (e.g., hardware fifo queue [13]) to improve the overall efficiency.

VII. Summary

TAPAS's primary goal is to provide an intuitive HLS toolchain for software programmers to generate parallel accelerators. We have decoupled concurrency from parallelism; we use the task-based programming framework to convey what can run in parallel and generate an architecture that can dynamically explore the available parallelism at run time. We hope this will be an effective framework for those in the community to build and exchange parallel accelerators. We have released *TAPAS* open source (github link redacted) which includes i) LLVM-based compiler back-end that translates parallel compiler IR to parallel accelerator architectures in Chisel, ii) a framework to convey concurrency and task parallelism to *TAPAS*, iii) Chisel libraries implementing support for task spawn/sync/reattach operations on an FPGA, iv) sample parallel accelerators (e.g., pipeline, nested loops, heterogeneous).

Acknowledgement

We would like to thank the anonymous reviewers and our shepherd for suggestions and feedback that helped to improve this paper.

References

- [1] Vivado Design Suite. <https://www.xilinx.com/products/design-tools/vivado.html>.
- [2] S N Agathos and V V Dimakopoulos. Compiler-Assisted OpenMP Runtime Organization for Embedded Multicores. *TR-2016-1, University of Ioannina*, 2016.

- [3] Ryo Asai and Andrey Vladimirov. Intel Cilk Plus for complex parallel algorithms - "Enormous Fast Fourier Transforms" (EFFT) library. *Journal of Parallel Computing*, 2015.
- [4] Jonathan Bachrach, Huy Vo, Brian Richards, Yunsup Lee, Andrew Waterman, Rimas Avizienis, John Wawrzyniek, and Krste Asanovic. Chisel: Constructing hardware in a scala embedded language. <https://github.com/freechipsproject/chisel3>.
- [5] David F Bacon, Rodric Rabbah, and Sunil Shukla. Fpga programming for the masses. *Communications of the ACM*, 56(4):56–63, 2013.
- [6] Lars Bauer, Artjom Grudnitsky, Muhammad Shafique 0001, and Jörg Henkel. PATS - A Performance Aware Task Scheduler for Runtime Reconfigurable Processors. In *Proc. of the FCCM*, 2012.
- [7] Robert D Blumofe and Charles E Leiserson. Scheduling Multithreaded Computations by Work Stealing. *Journal of ACM*, 1999.
- [8] Andrew Canis, Jongsok Choi, Blair Fort, Ruolong Lian, Qijing Huang, Nazanin Calagar, Marcel Gort, Jia Jun Qin, Mark Aldham, Tomasz Czajkowski, Stephen Brown, and Jason Anderson. From software to accelerators with LegUp high-level synthesis. In *In Proc. of CASES*, 2013.
- [9] George Charitopoulos, Iosif Koidis, Kyprianos Papadimitriou, and Dionisios Pnevmatikatos. Run-time management of systems with partially reconfigurable fpgas. *Integration, the VLSI Journal*, 57:34–44, 2017.
- [10] George Charitopoulos, Iosif Koidis, Kyprianos Papadimitriou, and Dionisios N Pnevmatikatos. Hardware Task Scheduling for Partially Reconfigurable FPGAs. In *Proc. of ARC*, 2015.
- [11] J. Choi, S. Brown, and J. Anderson. From software threads to parallel hardware in high-level synthesis for fpgas. In *Proc. of FPT*, 2013.
- [12] Jongsok Choi, Stephen Dean Brown, and Jason Helge Anderson. From pthreads to multicore hardware systems in legup high-level synthesis for fpgas. *IEEE Trans. VLSI Syst.*, 25(10), 2017.
- [13] Jongsok Choi, Ruolong Lian, Stephen Dean Brown, and Jason Helge Anderson. A unified software approach to specify pipeline and spatial parallelism in FPGA hardware. In *27th IEEE International Conference on Application-specific Systems, Architectures and Processors*, 2016.
- [14] Jongsok Choi, Kevin Nam, Andrew Canis, Jason Helge Anderson, Stephen Dean Brown, and Tomasz S Czajkowski. Impact of Cache Architecture and Interface on Performance and Area of FPGA-Based Processor/Parallel-Accelerator Systems. In *Proc. of the FCCM*, 2012.
- [15] Eric S Chung, James C Hoe, and Ken Mai. CoRAM: an in-fabric memory architecture for FPGA-based computing. In *Proc. of the 19th FPGA*, 2011.
- [16] J. Cong, B. Liu, S. Neuendorffer, J. Noguera, K. Vissers, and Z. Zhang. High-level synthesis for fpgas: From prototyping to deployment. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 30(4), 2011.
- [17] Mattias De Wael, Stefan Marr, and Tom Van Cutsem. Fork/join parallelism in the wild: Documenting patterns and anti-patterns in java programs using the fork/join framework. In *Proc. of the PPPJ*, 2014.
- [18] A DeHon, J Adams, M deLorimier, N Kapre, Y Matsuda, H Naeimi, M Vanier, and M Wrighton. Design patterns for reconfigurable computing. In *Proc. of the 12th FCCM*, 2004.
- [19] R Domingo, R Salvador, H Fabelo, D Madroñal, S Ortega, R Lazcano, E Juárez, G Callicó, and C Sanz. High-level design using intel fpga opencl: A hyperspectral imaging spatial-spectral classifier. In *Proc. of the Reconfigurable Communication-centric Systems-on-Chip (ReCoSoC)*, 2017.
- [20] Yoav Etsion, Felipe Cabarcas, Alejandro Rico, Alex Ramirez, Rosa M Badia, Eduard Ayguade, Jesus Labarta, and Mateo Valero. Task Superscalar: An Out-of-Order Task Pipeline. In *Proc. of the 43rd MICRO*, 2010.
- [21] Antonio Filgueras, Eduard Gil, Carlos Álvarez 0001, Daniel Jiménez-González, Xavier Martorell, Jan Langer, and Juanjo Noguera. Heterogeneous tasking on SMP/FPGA SoCs - The case of OmpSs and the Zynq. In *Proc. of VLSI-SoC*, 2013.
- [22] Matteo Frigo, Charles E Leiserson, and Keith H Randall. The Implementation of the Cilk-5 Multithreaded Language. In *Proc. of the PLDI*, 1998.
- [23] Anca Iordache, Guillaume Pierre, Peter Sanders, Jose Gabriel de F. Coutinho, and Mark Stillwell. High performance in the cloud with fpga groups. In *Proceedings of the 9th International Conference on Utility and Cloud Computing*, 2016.
- [24] Mark C Jeffrey, Suvinay Subramanian, Cong Yan, Joel S Emer, and Daniel Sanchez. A scalable architecture for ordered parallelism. In *Proc. of the 48th MICRO*, 2015.
- [25] Lana Josipović, Radhika Ghosal, and Paolo Ienne. Dynamically scheduled high-level synthesis. In *Proc. of the FPGA*, 2018.
- [26] N Kapre and H Patel. Applying Models of Computation to OpenCL Pipes for FPGA Computing. *Proc. of the 5th Intl. Workshop on OpenCL*, 2017.
- [27] Sanjeev Kumar, Christopher J Hughes, and Anthony Nguyen. Carbon: architectural support for fine-grained parallelism on chip multiprocessors. In *Proc. of the 34th ISCA*, 2007.
- [28] I-Ting Angelina Lee, Charles E. Leiserson, Tao B. Schardl, Zhunping Zhang, and Jim Sukha. On-the-fly pipeline parallelism. *ACM Transactions on Parallel Computing*, Oct 2015.
- [29] Zhaoshi Li, Leibo Liu, Yangdong Deng, Shouyi Yin, Yao Wang, and Shaojun Wei. Aggressive Pipelining of Irregular Applications on Reconfigurable Hardware. In *Proc. of ISCA*, 2017.
- [30] Yi Lu, Thomas Marconi, Koen Bertels, and Georgi Gaydadjiev. A Communication Aware Online Task Scheduling Algorithm for FPGA-Based Partially Reconfigurable Systems. *Proc. of the FCCM*, 2010.
- [31] Divya Mahajan, Jongse Park, Emmanuel Amaro, Hardik Sharma, Amir Yazdanbakhsh, Joon Kyung Kim, and Hadi Esmaeilzadeh. TABLA: A unified template-based framework for accelerating statistical machine learning. In *Proc. of the 22nd HPCA*, 2016.
- [32] Marc S Orr, Bradford M Beckmann, Steven K Reinhardt, and David A Wood. Fine-grain task aggregation and coordination on GPUs. In *Proc. of the 41st ISCA*, 2014.
- [33] Raghu Prabhakar, David Koeplinger, Kevin J Brown, HyounJoong Lee, Christopher De Sa, Christos Kozyrakis, and Kunle Olukotun. Generating Configurable Hardware from Parallel Patterns. In *Proc. of the 21st ASPLOS*, 2016.
- [34] Raghu Prabhakar, Yaqi Zhang, David Koeplinger, Matt Feldman, Tian Zhao, Stefan Hadjis, Ardavan Pedram, Christos Kozyrakis, and Kunle Olukotun. Plasticine: A reconfigurable architecture for parallel patterns. In *Proc. of the 44th ISCA*, 2017.
- [35] Andrew Putnam. FPGAs in the Datacenter - Combining the Worlds of Hardware and Software Development. *ACM Great Lakes Symposium on VLSI*, 2017.
- [36] Andrew Putnam, Dave Bennett, Eric Dellinger, Jeff Mason, and Prasanna Sundararajan. CHiMPS - a high-level compilation flow for hybrid CPU-FPGA architectures. In *Proc. of the FPGA*, page 261, 2008.
- [37] Andrew Putnam, Susan J Eggers, Dave Bennett, Eric Dellinger, Jeff Mason, Henry Styles, Prasanna Sundararajan, and Ralph Wittig. Performance and power of cache-based reconfigurable computing. *Proc. of the ISCA*, 2009.
- [38] Daniel Sanchez, Richard M Yoo, and Christos Kozyrakis. Flexible architectural support for fine-grain scheduling. In *Proc. of the 15th ASPLOS*, 2010.
- [39] Tao B Schardl, William S Moses, and Charles E Leiserson. Tapir - Embedding Fork-Join Parallelism into LLVM's Intermediate Representation. In *In Proc. of PPOPP*, 2017.
- [40] Shreesha Srinath, Berkin Ilbeyi, Mingxing Tan, Gai Liu, Zhiru Zhang, and Christopher Batten. Architectural Specialization for Inter-Iteration Loop Dependence Patterns. In *Proc. of the 47th MICRO*, 2014.
- [41] Olivier Tardieu, Haichuan Wang, and Haibo Lin. A work-stealing scheduler for X10's task parallelism with suspension. *Proc. of the 17th PPOPP*, 2017.
- [42] Hasitha Muthumala Waidyasooriya, Masanori Hariyama, and Kunio Uchiyama. FPGA-Oriented Parallel Programming. In *Design of FPGA-Based Computing Systems with OpenCL*. October 2017.
- [43] Christopher S Zakian, Timothy A K Zakian, Abhishek Kulkarni, Buddhika Chamith, and Ryan R Newton. Concurrent Cilk - Lazy Promotion from Tasks to Threads in C/C++. In *Proc. of LCPC*, 2015.